

Java #3

GC & Concurrent

Garbage Collection

Garbage Collector

- Memory Management Mechanism
- Used in most VM / interpreted environment
- Opposed to manual allocation

Garbage Collector

Pros:

- Extremely efficient
- Collection is free
- Death detection is precise

Cons:

- Stop 1s/Gb
- Does not prevent all memory leaks
- Pauses can be delayed, not stopped

Garbage Collector

Terminology

- **Monolithic** GC: operates in a single, indivisible step.
- **Parallel** GC: uses multiple threads to operate.
- **Concurrent** GC: operates while the application operates.
- **Stop-The-World** : pauses completely the application to collect.
- **Incremental** : operates in a sequence of divisible steps.
- Mostly: IS NOT.

Garbage Collector

Terminology

- **Conservative** : if the GC is unaware or unsure about some references while collecting.
- **Precise** : 100% accuracy.
- **GC safepoint** : point where the application can safely stop to let the GC operate.
- **Global Safepoint** : all threads are on a safe point -> Stop-The-World.

Garbage Collector

Phases

- Mark
- Sweep
- Compact
- Copy

All GC algorithms are combinations of these phases

Garbage Collector

Mark

- Start from roots :
 - Thread stacks,
 - Local Variables,
 - Statics;
- Mark anything reachable as "live".
- In the end, any non-marked object is dead.
- Complexity linear to "live" objects set, not heap size.

Garbage Collector

Sweep

- Scan through the heap.
- Sweep dead objects to a free list.
- Complexity linear to the heap size.

Garbage Collector

Compact

- Over time, fragmentation happens;
- Move live references to compact them (relocate),
- And reclaim consecutive free space.
- The hard part is updating references pointing to moved objects (remap)
- Complexity linear to live set.

Garbage Collector

Copy

- Copy all live objects from one space to another.
- Starting space is then completely free.
- Complexity linear to live set.
- Oftentimes, combined with Mark phase.

Garbage Collector

Generational Hypothesis

Most objects die young.

Garbage Collector

Generational Hypothesis

- Divide objects on generation:
 - Young (can be separated further),
 - Old.
- Promote objects that live long enough to an older gen.

Garbage Collector

Generational Hypothesis

- Collect young{er} generations with Mark/Copy, Monolithic STW.
- Collect old{er} generations as they fill up.
- Use Mark/Sweep for oldest gen.

Garbage Collector

Take away

- Young generation disposal is dirt cheap,
- Old generation is stable, but disposal is expensive.
- Design accordingly.

Concurrent Programming

Process:

- Self contained execution environment
- Complete, private set of resources, including memory space
- Communication between process through IPCs:
 - Pipes
 - Sockets
 - Shared Memory

Threads:

- Lightweight process
- Exist within a process
- Every process has at least one thread
- Process resources are shared between threads

Threads:

- Manipulating threads is hard:
 - Synchronization
 - Data access

Alternatives:

- Actors
- Event Loops
- Executors
- Parallel Streams
- Fork/Join

Synchronizing data accesses

- Concurrent data accesses are unsafe
- Most operations are not atomic
- even a simple ++i may cause issue
- Some problems are more complex

Synchronizing data accesses

Lock/Mutex

- Monitor:
 - Object protected with a lock (mutex)
 - Invoking methods locks it
- All java object implements a monitor
- `synchronized` method are monitor methods
- Sufficient for all basic situations
- Locking is reentrant:
 - You can re-acquire a lock if you hold it already
 - Call to other synchronised methods are safe

Synchronizing data accesses

Lock/Mutex

```
public class SimpleSyncCounter {  
    private int counter;  
  
    public SimpleSyncCounter() { counter = 0; }  
  
    public synchronized void increment() {counter += 1;}  
    public synchronized void decrement() {counter -= 1;}  
    public synchronized int get() {return counter;}  
}
```

Synchronizing data accesses

Lock/Mutex

```
public class SimpleAtomicCounter {  
    private AtomicInteger counter = new AtomicInteger(0);  
    public int get() { return counter.get(); }  
    public void increment() { counter.getAndIncrement(); }  
    public void decrement() { counter.getAndDecrement(); }  
}
```


Synchronizing data accesses

Atomic Types

- Useful basic types in atomic version:
 - int, boolean, int[], reference ...
 - Note that double and long access is not atomic.
- More efficient than locked methods
- Offers usual atomic operations
- Needed to implement efficient concurrent data structures (non-blocking algorithms).

Synchronizing data access

Atomic Types

- All members marked as volatile can be accessed atomically.
- Any change to a volatile property creates a "happened-before" relationship.
 - IE: The vm "commits" the change to make sure it is immediately visible to all threads.

Synchronizing data access

Synchronized Blocks

```
public class IncCounter {  
    private Integer c1 = 42;  
    public void add(final Integer toAdd) {  
        // More control over when synchronization happens.  
        if (toAdd > 0) {  
            synchronized (c1) {  
                c1 += toAdd;  
            }  
        }  
    }  
}
```

Synchronizing data accesses

Synchronized Blocks

```
// Finer access control allows better liveness.
public class DualQueue<T> {

    private final Queue<T> firstQueue = new ArrayDeque<>();
    private final Queue<T> secondQueue = new ArrayDeque<>();

    public void pushFirst(final T element) {
        // Other threads can still access to secondQueue
        synchronized (firstQueue) {
            firstQueue.add(element);
        }
    }

    public void pushSecond(final T element) {
        synchronized (secondQueue) {
            // Other threads can still access to firstQueue
            secondQueue.add(element);
        }
    }
}
```

Synchronizing data access

Synchronized Blocks

- `Object#wait`
 - Current thread wait, with the object as the monitor.
- `Object#notify`
 - Wakes one of the threads waiting on the monitor at random.
- `Object#notifyAll`
 - Wakes all the threads waiting on the monitor.

Blocking Queues

- Efficient scalable thread-safe non-blocking FIFO queue
- Async put/take if not full, can wait or throw otherwise
- Excellent for communication between threads

Blocking Queues

```
class BlockingQueueExample {  
    public static void main(String[] args) throws Exception {  
        BlockingQueue queue = new ArrayBlockingQueue(1024);  
        Producer producer = new Producer(queue); Consumer consumer = new Consumer(queue);  
        new Thread(producer).start();  
        new Thread(consumer).start();  
        Thread.sleep(4000);  
    }  
}
```

CompletableFuture

- Interface
- Most modern tool for concurrent development in the JDK
- Defines the contract for an asynchronous computation step that can be combined with other steps.
- Execution can be delegated to a scheduler thread pool
- Manage exceptions gracefully

CompletableFuture

```
class TotallyNotUsefulDuringTheRush {  
    public static String doSomethingWithAFixedDelay() {...}  
    public static void main(String[] args) throws Exception {  
        CompletableFuture  
            .supplyAsync(() -> doSomethingWithAFixedDelay())  
            .get();  
    }  
}
```

CompletableFuture

```
import java.util.concurrent.CompletableFuture;

class TotallyNotUsefulDuringTheRush {
    public static String doSomethingWithAFixedDelay() {...}

    public static String fetchReport(final String reportId) {...}

    public static void main(String[] args) throws Exception {
        CompletableFuture.supplyAsync(() -> doSomethingWithAFixedDelay())
            .thenApplyAsync(reportId -> fetchReport(reportId))
            .get();
    }
}
```

CompletableFuture

```
class TotallyNotUsefulDuringTheRush {  
    public static String doSomethingWithAFixedDelay() {...}  
    public static String fetchReport(final String reportId) {...}  
    public static void main(String[] args) throws Exception {  
        final long delay = 10L;  
        final Executor executor =  
            CompletableFuture.delayedExecutor(delay, TimeUnit.MILLISECONDS);  
        CompletableFuture  
            .supplyAsync(() -> doSomethingWithAFixedDelay())  
            .thenApplyAsync(reportId -> fetchReport(reportId), executor)  
            .get();  
    }  
}
```

CompletableFuture

Recovering from errors

```
class TotallyNotUsefulDuringTheRush {  
    public static String doSomethingThatMayFail() {...}  
  
    public static void main(String[] args) throws Exception {  
        final long delay = 10L;  
  
        final Executor executor =  
            CompletableFuture.delayedExecutor(delay, TimeUnit.MILLISECONDS);  
  
        CompletableFuture  
            .supplyAsync(() -> doSomethingThatMayFail())  
            .handle((value, exception) -> return value != null ? value : "ERROR")  
            .get();  
    }  
}
```

CompletableFuture

Exercise

- Write a sequence of completable future that:
 - Produces a random int between 0 and 10 included.
 - Wait one second.
 - Divide the generated value by another random value between -1 and 1.
 - Prints the result.
 - Errors should be recovered by generating the value 42.